

TensorFlow Dropout




Figure 1: Taken from the paper "Dropout: A Simple Way to Prevent Neural Networks from Overfitting" (<https://www.cs.toronto.edu/~hinton/absps/JMLRdropout.pdf>)

Dropout is a regularization technique for reducing overfitting. The technique temporarily drops units (**artificial neurons**) from the network, along with all of those units' incoming and outgoing connections. Figure 1 illustrates how dropout works.

TensorFlow provides the **`tf.nn.dropout()`** function, which you can use to implement dropout.

Let's look at an example of how to use **`tf.nn.dropout()`**.

```
keep_prob = tf.placeholder(tf.float32) # probability to keep units
```

```
hidden_layer = tf.add(tf.matmul(features, weights[0]), biases[0])  
hidden_layer = tf.nn.relu(hidden_layer)  
hidden_layer = tf.nn.dropout(hidden_layer, keep_prob)
```

```
logits = tf.add(tf.matmul(hidden_layer, weights[1]), biases[1])
```

The code above illustrates how to apply dropout to a neural network.

The `tf.nn.dropout()` function takes in two parameters:

1. `hidden_layer`: the tensor to which you would like to apply dropout
2. `keep_prob`: the probability of keeping (i.e. *not*dropping) any given unit

`keep_prob` allows you to adjust the number of units to drop. In order to compensate for dropped units, `tf.nn.dropout()` multiplies all units that are kept (i.e. *not*dropped) by $1/\text{keep_prob}$.

During training, a good starting value for `keep_prob` is `0.5`.

During testing, use a `keep_prob` value of `1.0` to keep all units and maximize the power of the model.

Quiz 1

Take a look at the code snippet below. Do you see what's wrong?

There's nothing wrong with the syntax, however the test accuracy is extremely low.

...

```
keep_prob = tf.placeholder(tf.float32) # probability to keep units
```

```
hidden_layer = tf.add(tf.matmul(features, weights[0]), biases[0])
```

```
hidden_layer = tf.nn.relu(hidden_layer)
```

```
hidden_layer = tf.nn.dropout(hidden_layer, keep_prob)
```

```
logits = tf.add(tf.matmul(hidden_layer, weights[1]), biases[1])
```

...

```
with tf.Session() as sess:
```

```
    sess.run(tf.global_variables_initializer())
```

```
    for epoch_i in range(epochs):
```

```
        for batch_i in range(batches):
```

```
            ....
```

```
            sess.run(optimizer, feed_dict={
```

```
features: batch_features,  
labels: batch_labels,  
keep_prob: 0.5})
```

```
validation_accuracy = sess.run(accuracy, feed_dict={  
    features: test_features,  
    labels: test_labels,  
    keep_prob: 0.5})
```

QUESTION 1 OF 2

What's wrong with the above code?

Dropout doesn't work with batching.

The keep_prob value of 0.5 is too low.

There shouldn't be a value passed to keep_prob when testing for accuracy.

keep_prob should be set to 1.0 when evaluating validation accuracy.

SUBMIT

Quiz 2

This quiz will be starting with the code from the ReLU Quiz and applying a dropout layer. Build a model with a ReLU layer and dropout layer using the **keep_prob** placeholder to pass in a probability of **0.5**. Print the logits from the model.

Note: Output will be different every time the code is run. This is caused by dropout randomizing the units it drops.